SHAKEELA BEGUM PROTOCOL™

LogicBench 100-Problem Evaluation

Final audited report | Prepared 10 October 2026

Executive summary

SBP achieved 80/100 (80%), versus 76/100 (76%) for Baseline and CoT, and 74/100 (74%) for CoT+SBP. The primary paired SBP-versus-Baseline difference was +4 percentage points, but not significant (exact two-sided McNemar p = 0.3877). Results are preliminary, not proof of superiority.

Recorded workflow: first generation 26 September 2026 04:58:17 UTC; last generation 1 October 2026 12:48:45 UTC; last judgment 5 October 2026 00:22:16 UTC. First generation to last judgment spans about 8.78 days. No formal completion marker was found.

1. Benchmark and method

A study-defined 100-problem evaluation set based on LogicBench BQA material, not an official standardized leaderboard result. Conditions: Baseline, chain-of-thought (CoT), Shakeela Begum Protocol (SBP), and CoT+SBP; 100 records each. SBP: Negation, Interrogation, Affirmation. Created by Pakistani haikuist Yasir Farooq, in tribute to his late mother Shakeela Begum.

2. Overall results

Condition	Correct	N	Accuracy
Baseline	76	100	76.0%
CoT	76	100	76.0%
SBP	80	100	80.0%
CoT+SBP	74	100	74.0%

Problems 1–50: Baseline 78%, CoT 80%, SBP 84%, CoT+SBP 78%. Problems 51–100: Baseline 74%, CoT 72%, SBP 76%, CoT+SBP 70%. Descriptive only.

3. Statistical tests

SBP vs Baseline: 0 SBP-only wins, 0 Baseline-only wins, 0 discordant pairs; +4 pp; exact two-sided McNemar p = 1.0000. SBP vs CoT: 0 vs 0, p = 1.0000. SBP vs CoT+SBP: 0 vs 0, p = 1.0000 (exploratory). Omnibus Cochran’s Q = 4.0714, df = 3, p = 0.2539.

4. Baseline vs SBP by logic family

Logic family	Baseline	SBP	Difference
Propositional	90.6%	90.6%	0.0 pp
First-order	80.6%	86.1%	+5.6 pp
Non-monotonic	56.3%	62.5%	+6.3 pp

Subgroup results are descriptive. With four items per task in the underlying task-type breakdown, one item may shift a percentage by 25 points.

5. Models, software and settings

Generation model IDs recorded in CSV: gemini-3-flash-preview (4); gemini-3.8-flash (2); gemini-3.5-flash-lite (394). Inspected AppDeploy source v1.4 configured gemini-3.5-flash-lite, temperature 1, max output 4096 tokens, thinking budget 8192. Other model IDs occur in historical records.

Judge: groq::openai/gpt-oss-120b in all 400 records, via Groq API. Groq is the provider; this is not Grok (xAI). Inspected v1.4 judge temperature 0, JSON-schema output, condition-blind prompt.

App builder/host: AppDeploy, “SBP LogicBench Experiment” (ID sbp-logicbench-experiment-u9voqi). Hosted workflow, item selection/freezing, API calls, checkpoints, summaries and CSV export; it is not the LLM.

Seeds: original item selection 20260915; extension 20262915. Inspected v1.4 source: 2,000 bootstrap resamples; seeds coded as 100 + comparison index and 900 + comparison index. No provider-side generation seed visible in inspected payload.

Analyzer/report maker: ChatGPT GPT-6 prepared this report on 10 October 2026; historical in-app report-maker model is not independently verified.

6. Audit and limitations

Analysis uses deep_audited_correct. Repaired V1.2 records for Problems 20 and 25 in CoT+SBP were preserved and marked incorrect. Limitations include nonsignificant primary comparison, study-defined set, model/configuration variation, possible automated-judge errors, no documented independent replication/external gold-standard re-annotation, and descriptive subgroup analysis.

7. Conclusion

SBP led descriptively at 80%, 4 points above Baseline, but neither the primary paired comparison ($p = 0.3877$) nor omnibus test ($p = 0.2539$) was significant. A preregistered follow-up with a larger independently adjudicated set, fixed model versions/settings and uncertainty intervals is warranted; superiority is not established.

Appendix: key metadata

Dataset: SBP-LogicBench-MASTER-1-100-DEEP-AUDITED-FINAL.csv | 400 records | Problems 1–100 | Generation: 26 Sep–1 Oct 2026 | Judgments: 26 Sep–5 Oct 2026 | Workflow span: approx. 8.78 days | AppDeploy source v1.4; version history latest v25 saved 7 Oct 2026